

Container Virtualization

Docker by Example

September 11, 2018

Marcel Großmann, Stefan Kolb, Dr. Andreas Schönberger, Gabriel Nikol

Content

1 Docker Concepts

2 Docker CLI

3 Dockerfile

4 Docker Compose



Docker Concepts

Docker CLI

Dockerfile

Docker Compose

What you can expect (and what not)



Docker Concepts

Docker CLI

Dockerfile

Docker Compose

What you learn

- you will learn the concept of docker
- you will learn the most used commands (not representative)
- you get a toolset to dockerize your applications
- you get some practice with the tools (use the tools in exercises)

But: Everything you will learn is in the docker documentation, too.

- you won't be a pro with docker (this needs practice)
- if you only dockerize your application from time to time this is mostly enough
- sometime you need to look into the documentation (or stackoverflow) for special cases/challenges
- this is not a guide to configure proxies, webserver etc. we focus on docker

What does Docker do:

- “Docker allows you to **package an application with all of its dependencies into a standardized** unit for software development.”
- “Docker **containers** wrap up a piece of software in a complete filesystem that **contains everything it needs to run:** code, runtime, system tools, system libraries - anything you can install on a server. This guarantess that it will always run the same, regardless of the environment it is running in.”

These artifacts are packed and can be exchanged.

- **Build, ship and run everywhere!**

Background: Evolution of Virtualization

- Lightweight containers enable easy horizontal scalability.

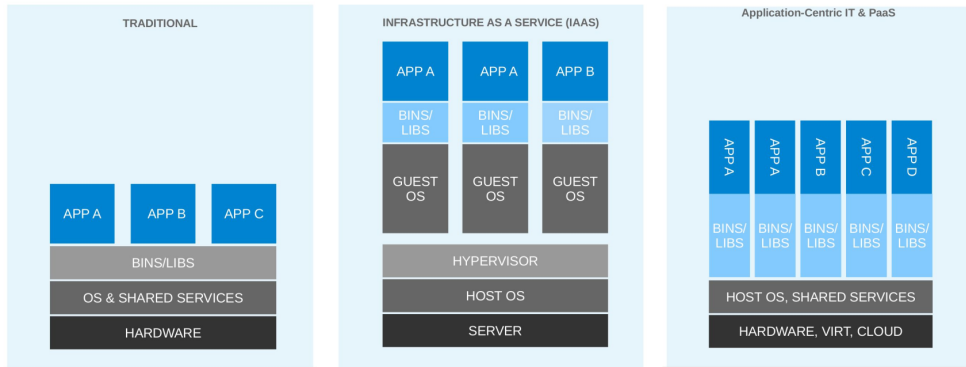


Figure 1: Evolution of Virtualization

Hypervisor

- Runs Operating System
- Heavyweight isolated virtual machines
- Can theoretically emulate any architecture
- VMs start via a full boot-up process
- Platform-oriented solution
- Optimized for generality

Container-enabled Kernel

- Runs processes
- Lightweight kernel namespaces
- Is less flexible in architecture emulation
- Very fast namespace + process creation
- Service-oriented solution
- Optimized for minimalism and speed

Figure 2: Virtual Machines vs. Container

Docker's Execution Drivers

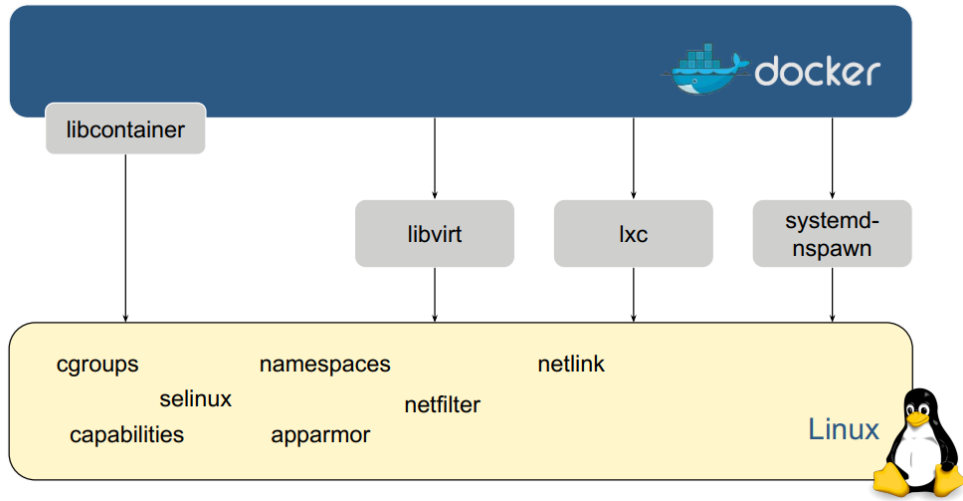


Figure 3: Docker's Execution Drivers

Container vs. Images

- An image is blueprint for a container
- An image contains of a read-only set of filesystem layers
- An instance of an image is called container
- You can have many running containers of the same image
- Major difference: container adds a writeable filesystem layer on top of the image




Figure 4: Containers vs Images

Changes and Updates

Copying makes containers eddicient

- only filesystem deltas need to be saved
- multiple containers can safely share a single underlying image

Changes made to a running container can always be committed to a new image




Figure 5: Changes And Updates

Data Volumes

- When a container is deleted, any data written to the container that is not stored in a *data volume* is deleted along with the container.
- A data volume is a directory or file in the host's filesystem that is mounted directly into a container
- Multiple volumes can be mounted by a container and volumes can also be shared between containers

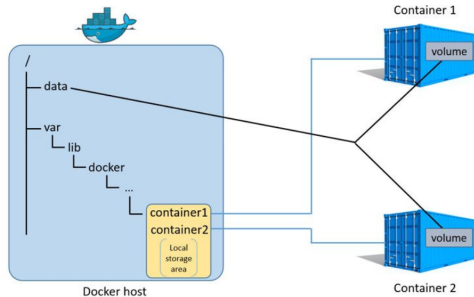


Figure 6: Data Volumes

Overall Architecture

- Docker CLI (Command Line Interface) is the docker client




Figure 7: Docker overview

- **Daemon:** daemon of the server process to manage containers
- **Client:** user client to (remotely) control the daemon
- **Registry:** platform for sharing and managing images

Shows **state** of containers, volumes and images.

Show **running** container only:

- *docker container ls*

Show **all** containers (both commands do the same):

- *docker container ls -a*
- *docker container ls --all*

The same for volumes, networks and images:

- *docker volume ls --all*
- *docker image ls --all*
- *docker network ls*

The **docker pull IMAGE** command downloads an image from docker registry.

In our case (default) from the **public** registry - docker hub

Pull the **latest** image of busybox (default):

- `docker pull busybox`

Pull busybox with tag **1.29.2-glibc**:

- `docker pull busybox:1.29.2-glibc`

The **docker run [OPTIONS] IMAGE[:TAG]** command starts container from images.

- `docker run busybox`

Start container for interactive processes (*-i*) and allocate a tty (*-t*). Together written as *-it*.

Open busybox with shell:

- `docker run -it busybox`

Other very important docker run **[OPTIONS]**: Give container names (unique):

- `docker run -it --name busy busybox`

delete container foobar after running:

- `docker run --rm busybox`

You can also inspect container, volumes, networks and images giving you detailed information:

- `docker container inspect ID|NAME`
- `docker volume inspect VOLUMENAME`
- `docker image inspect IMAGE`
- `docker network inspect ID|NAME`

Exercise 1

- 1 Open a first (power)shell.
- 2 download the image: `busybox` and `ubuntu`.
What do you see? What is different between the two images?
- 3 start a **ubuntu** container with the name **foo**.
List all directories (`ls`), create a new directory (`mkdir DIRNAME`), List all directories again.
- 4 Open a second (power)shell.
- 5 List all containers. What is the Id and name of your container?
- 6 Go back to the first (power)shell and exit the container.
Grateful with `exit`, abort with `STRG+C`.
- 7 List **all** containers again. What changes?
- 8 inspect the ubuntu container and image
- 9 delete the container by using: `docker container rm ID|NAME` and list all containers again



Docker Concepts

Docker CLI

Dockerfile

Docker Compose

Docker handling detached containers

Normally container run in the background without interaction.

Run container without active tty with *-detach* or *-d* option:

- `docker run -detach -name influx -rm influxdb`

Inspect stdout with *logs*

- `docker logs influx`

Stop running container with

- `docker stop influx`

Stop running container with

- `docker stop influx`



Docker Concepts

Docker CLI

Dockerfile

Docker Compose

You can define environment variables passed to a container.

Syntax: `-e=VARIABLE` or `--env=VARIABLE`

Example: print out value of environment variable foo on console

- `docker run -it -e foo=bar ubuntu`
- `/ # echo $foo`

publish ports



Docker Concepts

Docker CLI

Dockerfile

Docker Compose

By default no ports are accessible outside of container networks. You need to make the explicitly available to the host.

Syntax: **-p EXTERN:INTERN** or **-publish EXTERN:INTERN**

Example: start a http server on port 80

- `docker run -p 80:80 httpd`
- `docker run -publish 80:80 httpd`

Mapping of local file system into container file system.

Syntax:

- **-v LOCALFS:CONTAINERFS**
- **- -volume LOCALFS:CONTAINERFS**

Local file system path: C:/User/Name/html Example: start php container with C:/User/Name/html mapped into /var/www/html

- `docker run -v C:/User/Name/html:/var/www/html php:apache`

Exercise 2

- 1 GoTo: https://hub.docker.com/_/mysql/ and find the setable environment variables.
- 2 Run all container in detached mode
- 3 Run a mysql container named database with root password **foobar**.
- 4 Check the logs of the mysql
- 5 Stop the mysql container
- 6 Run a php:7.2-apache container named webserver, publish port 80 and map directory Exercise2/html into /var/www/html
- 7 Ensure that the webservice's html page is available
- 8 list all running container and compare the output with Figure 8. It should be equal.



Figure 8: *Exercise2: docker container ls*

Volumes without mounted host directory, so no absolute path needed. Increases portability of applications.

Create a volume named database.

- `docker volume create database`

Remember: list volumes

- `docker volume ls`

Map a volume like a file system volume by its name.

- `docker run -v database:/var/lib/influxdb --name influx --rm influxdb`

Exercise 3



Docker Concepts

Docker CLI

Dockerfile

Docker Compose

- 1 open a first (power)shell
- 2 create a volume named `shared_volume`
- 3 run a ubuntu image and map `shared_volume` to `/shared_volume`
- 4 open a second (power)shell
- 5 run a influxdb image and map `shared_volume` to `/var/lib/influxdb`
- 6 change back to the first (power)shell
- 7 list (`ls`) the content of `/shared_volume`

Until now, containers were not interconnected (except over published ports). We can link containers via their names, this is provided by an internal Docker DNS server. For this, they must be in the same Docker network.

Create a network named mynet.

- `docker network create mynet`

Remember: list networks

- `docker networks ls`

Join a container to a network:

- `docker run --name somecontainer --net mynet busybox`

Exercise 4



Docker Concepts

Docker CLI

Dockerfile

Docker Compose

- 1 open a first (power)shell
- 2 create a network named foobar
- 3 run a busybox image named foo and join the foobar
- 4 open a second (power)shell
- 5 run a busybox image named bar and join the foobar
- 6 ping foo from the second and bar from the first (power)shell
- 7 stop container foo, what happens in the bar container?
- 8 inspect the foobar network
- 9 stop all containers, delete the network and clean up the environment

You can prune (delete unused) containers, volumes, networks, images.

- docker container prune
- docker image prune
- docker volume prune
- docker network prune

You can do all at once with:

- docker system prune

Afterwards check with docker [container | volume | network | image] ls

Other usefull commands:

Execute something inside a running container - in most cases a shell: Example: detached httpd webserver and exec into bash

- `docker run -d -p 80:80 --name webserver httpd`
- `sudo docker exec -i -t webserver /bin/bash`

Override default entrypoint with docker run: Example: instead of executing apache go to bash

- `docker run -it --entrypoint "/bin/bash" php`

Motivation

- until now we used existing images
- you can commit changes to created containers, this creates new images (not addressed in this tutorial)
- but cumbersome to share local created images, hard to reproduce
- solution: declarative way to create images => Dockerfiles

Dockerfile - Concept

Concept

- blueprints for images
- a list of commands to create a state of an image
- makes image creation reproducible
- each RUN instruction adds a delta-layer

Instructions

- single line statements and contain a keyword
- not case sensitive

Example:



Figure 9: *Minimalistic Dockerfile*



Docker Concepts

Docker CLI

Dockerfile

Docker Compose

Every Dockerfile starts with an image it is based on. For this the FROM Keyword is used.

- FROM <image>[:<tag>] [AS <name>]

Example: Ubuntu is a good start for every new docker application

- FROM ubuntu:latest

You can add comments to you Dockerfile with the # (hash) Example: A comment

- # this is comment in a Dockerfile

A Dockerfile is a list of instructions to create a image. To execute commands on the shell use the RUN keyword. There exist two forms (shell and exec form): shell form:

- RUN <command>

exec form (this form is required for paths containing whitespace):

- RUN ["executable", "param1", "param2"]

In the shell form you can use a \ (backslash) to continue a single RUN instruction onto the next line. This increases readability!

Example: install git inside an ubuntu image and clean up apt list for thin layers

- RUN apt-get update && apt-get install -y \
git \
&& rm -rf /var/lib/apt/lists/*

Until now you never explicitly said what application to run when you started a container. When writing dockerfile you need to define this by using the `CMD` or `ENTRYPOINT` command. This sets the image's main command.

`CMD`

- `CMD ["executable", "param1", "param2"...]`
- `CMD executable param1`

`ENTRYPOINT` (exec form preferred)

- `ENTRYPOINT ["executable", "param1", "param2"...]`
- `ENTRYPOINT executable param1`

You can build your images from a Dockerfile by using the Docker CLI.

1 path to Dockerfile (Default is 'PATH/Dockerfile') Syntax:

- -f path/to/Dockerfile
- --file path/to/Dockerfile

2 name and tag of image Syntax:

- -t name:tag
- --tag name:tag

3 last option is the path to the build context (mostly .)

Example: Build a Dockerfile with buildcontext in current (power)shell directory

- `docker build -t foo:bar -f path/to/Dockerfile .`

Dockerfile - Exercise 1

- 1 go to directory `dockerfile/Exercise1`
- 2 create a new Dockerfile and use `ubuntu` as base image
- 3 install `cowsay` and `lolcat` in one layer (optional: clean up)
- 4 create an ENTRYPOINT in exec format executing `/usr/games/cowsay` and *"Hello"* as first parameter
- 5 create a CMD in exec format with *"World"* (will be second parameter to `cowsay`)
- 6 build your image named `cowsay`
- 7 run the image `cowsay` without parameter and with parameter (do not change the Dockerfile)
- 8 comment out the ENTRYPOINT and CMD
- 9 append a new CMD with shell format executing the following:
`/usr/games/cowsay "What does the cow say?" | /usr/games/lolcat`
- 10 build your image named `lolsay`
- 11 run the image `lolsay`



Docker Concepts

Docker CLI

Dockerfile

Docker Compose

Until now you never added local files like prebuild executables or html/php sides to the images. Each of these commands adds a new delta-layer to the images. You can add files by using ADD and COPY.

Syntax (chown only works on linux containers):

- ADD [-chown=<user>:<group>] <src>... <dest>
- ADD [-chown=<user>:<group>] ["<src>",... "<dest>"]

Example: adds all files starting with "hom"

- ADD files* /somedir/

Syntax (chown only works on linux containers):

- COPY [-chown=<user>:<group>] <src>... <dest>
- COPY [-chown=<user>:<group>] ["<src>",... "<dest>"]

Example: ? is replaced with any single character, e.g., "home.txt"

- COPY hom?.txt /mydir/

You need to explicitly tell docker in your dockerfiles which port are used. Only these ports can be published during docker run -p EXTERNALPORT:INTERNALPORT
Syntax:

- EXPOSE INTERNALPORT

Example: expose port 8081 and publish it to 8080

1 in dockerfile:

- EXPOSE 8081

2 docker run command:

- docker run -p 8080:8081 imagename

Dockerfile - Exercise 2

- go to directory `dockerfile/Exercise2`
- create a new Dockerfile
- use `ubuntu` as base image
- install `apache2` (and cleanup)
- add `html` to `/var/www/html`
- copy `etc/apache2/ports.conf` to `/etc/apache2/ports.conf`
- copy `etc/apache2/site-available/000-default.conf` to `etc/apache2/site-available/000-default.conf`
- expose port `8080`
- set entrypoint to start `apache2` with `apachectl -D FOREGROUND`
- create image named `ubuntu` with the tag `apache2`
- run your created image, publish port to `80` (optional run it in detached mode)
- check if your webside is locally reachable



Docker Concepts

Docker CLI

Dockerfile

Docker Compose

There are many more Dockerfile Instruction e.g.

- WORKDIR PATH
- ENV KEY VALUE
- ARGS
- VOLUME creates a volume at runtime
- many more. e.g.

Docker Compose - Motivation

- Typically, an application is not one service, but orchestration of multiple smaller isolated service units




Figure 10: *NGinx, PHP, MySQL Service-Stack*

Multi-container apps are a hassle

- Until now:
- pull images, build dockerfiles
- create networks
- create container with long configuration lines (volumes, publish, networks)
- naming is global
- container logs for each container by name or id
- for a lot of containers!!

Multi-container apps are a hassle

- docker-compose up does everything automatically (after configuration)

Docker Compose File - Overview

- must be named docker-compose.yml or docker-compose.yaml
- containers become services
- yaml file (use no tabulator, space only)

```
1  version: '3'
2
3  services:
4    httpserver:
5      image: httpd
6      ports:
7        - 80:80
8      volumes:
9        - ../code
10       - logvolume01:/var/log
11
12    customservice:
13      build:
14        context: .
15        dockerfile: Dockerfile
16      environment:
```



Docker Concepts

Docker CLI

Dockerfile

Docker Compose

There exists a wide range of docker-compose cli commands. See docker-compose cli overview. Most important commands:

start services in foreground:

- `docker-compose up`

start services detached:

- `docker-compose up -d`

show service(s) log (console output):

- `docker-compose logs [servicename]`

stop services (force deletes running container)

- `docker-compose rm [-force] | [-stop]`

Docker Compose File - Ports

publish becomes ports with a list of port mappings.

```
3  services:
4      httpserver:
5          image: httpd
6          ports:
7              - 80:80
```

Figure 12: Example: publish port 80 to 80

Docker Compose File - Volumes

Local volume mapping stays volumes with a list of volume mappings. Volumes are listed at the end of the yml-file

Example: local mapping and with volume logvolume01

```
8      volumes:
9      - ./code
10     - logvolume01:/var/log
```

Figure 13: Example: Mapping of . to /db, logvolume01 to /var/log

Example: volume definition of logvolume01

```
23     volumes:
24     logvolume01:
```

Figure 14: Example: volumes definitions

Docker Compose File - Build Images with Docker Compose



Docker Concepts

Docker CLI

Dockerfile

Docker Compose

Add a build command inside service. Image's name and tag is generated automatically by the service's name. The context and path to the dockerfile can be set by key-value

```
12     customservice:  
13         build:  
14             context: .  
15             dockerfile: Dockerfile
```

Figure 15: *Example: build with context and dockerfile*

Create the images then with:

- docker-compose build

Docker Compose File - Environment Variables



Docker Concepts

Docker CLI

Dockerfile

Docker Compose

Add a environment command inside service with a list of environment key-values
Example: `docker run -e Foo=Bar ...`

```
16      environment:  
17      - Foo=Bar
```

Figure 16: *Example: environment variable key=foo and value=bar*

Exercise Docker Compose - PHP Application 1/2



Docker Concepts

Docker CLI

Dockerfile

Docker Compose

First Service: PHP with Apache2

- Dockerfile
 - expose port 8080
 - create dockerfile from php:apache2
 - copy scripts to /var/www/html
- publish port 8080
- create a new service named phpapache2

Second Service: Database with MySQL

- create a new service named db_mysql
- map volume
- define environment variables user passwd

Docker Compose File - Command Translation

Add the networks: keyword at the end of the yml-file. Add a list of networks.

Example: Usage of networks in service

```
4      httpserver:
5        image: httpd
6        ports:
7          - 80:80
8        volumes:
9          - ./code
10         - logvolume01:/var/log
11        networks:
12          - example_network
13
```

Figure 17: Example: usage of network `example_network` in service `httpserver`

Example: network definition of `example_network`

```
26    networks:
27      example_network:
```

Figure 18: Example: networks definition of `example_network`

Exercise Docker Compose - PHP Application with NGINX 2/2



Docker Concepts

Docker CLI

Dockerfile

Docker Compose

User your artifact from the last exercise First Service: NGINX

- create a new service named `nginx_reverse_proxy`
- expose port 80
- volume `conf.d/default.conf` into
- depend on `php_apache2`
- join an appropriate network

Edit Service: PHP with Apache2

- remove the publish

Second Service: Database with MySQL

- create a new service named `db_mysql`
- map volume etc

Networks

- the Database service shall only be reachable by the `php_apache2` service, so put them in the same network
- the `php_apache2` shall be reachable from the `nginx` reverse proxy but not from the outside
- `nginx` shall be accessible over port 80 from the outside

Questions ?

made by Gabriel Nikol

<https://github.com/S3ler/DockerByExample>